

14/8/26

Week-4

Aim: To build a simple feed forward Neural Network to recognize handwritten character, using MNIST dataset.

Objective:-

- 1) To understand the architecture of a feed forward Neural Network.
- 2) To Apply supervised learning for image classification.
- 3) To preprocess MNIST dataset for Neural Network training.
- 4) To evaluate the models accuracy and analyze the performance.

Procedure:-

- 1) Import libraries (Tensorflow, keras, matplotlib)
- 2) Load MNIST dataset.
- 3) Normalize the data (pixel values)
- 4) one-hot encode labels (0-9 into vector length of 10)
- 5) Define model architecture:
 - Flatten layer (convert 28×28 images into 784 inputs)
 - Dense hidden layer with ReLU activation.
 - Dense output layer with softmax activation
- 6) Compile model with optimizer, loss function, metrics.
- 7) Train the model on training data.
- 8) Evaluate performance on test data.
- 9) plot accuracy graph for train and validation sets.

Pseudo code:-

START

Import Tensorflow, keras, matplotlib

Load MNIST dataset

Normalize Input Images to range $[0, 1]$

Convert labels to one-hot encoding

Initialize sequential model

Add flatten layer for 28×28 input

Add dense hidden layer with ReLU activation
Add dense output layer with softmax activation
compile model with Adam optimizer, categorical cross entropy loss, accuracy metrics

train model for 5 epochs

Evaluate model on test data

plot training vs validation accuracy

predict labels for first 5 test samples

display predicted labels with images

END

Observations:-

- *) The model achieved $\sim 97-98\%$ accuracy on the MNIST test dataset with just 5 epochs of training.
- *) Accuracy improved steadily with each epoch, indicating effective learning.
- *) predictions on unseen data matched the actual digits in most cases.
- *) Errors occurred mainly on digits that were poorly written.

conclusion:-

A simple feed forward neural network with one hidden layer can accurately classify handwritten digits from the MNIST dataset.

The performance shows that FFNN's are effective for basic image recognition tasks after normalization and one-hot encoding.

However, for more complex image recognition problems, deeper architectures like convolutional neural networks may be more suitable.

output :-

Epoch 1/5

val_accuracy : 0.9596 - val_loss : 0.1386

Epoch 2/5

val_accuracy : 0.9706 - val_loss : 0.0983

Epoch 3/5

val_accuracy : 0.9716 - val_loss : 0.0882

Epoch 4/5


val_accuracy : 0.9759 - val_loss : 0.0763

Epoch 5/5


val_accuracy : 0.9748 - val_loss : 0.0815

Test accuracy : 0.9748


predicted : 7




predicted : 0




predicted : 2



predicted : 4



predicted : 1



Q.1

```
import tensorflow as tf
from tensorflow.keras.datasets import mnist
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Flatten
from tensorflow.keras.utils import to_categorical
import matplotlib.pyplot as plt

(x_train, y_train), (x_test, y_test) = mnist.load_data()
x_train = x_train / 255.0
x_test = x_test / 255.0
y_train = to_categorical(y_train, 10)
y_test = to_categorical(y_test, 10)
model = Sequential()
model.add(Flatten(input_shape=(28, 28)))
model.add(Dense(128, activation='relu'))
model.add(Dense(10, activation='softmax'))
model.compile(optimizer='adam',
              loss='categorical_crossentropy',
              metrics=['accuracy'])
history = model.fit(x_train, y_train,
                   epochs=5,
                   batch_size=32,
                   validation_data=(x_test, y_test))
test_loss, test_acc = model.evaluate(x_test, y_test)
print(f"\n Test Accuracy: {test_acc:.4f}")
plt.plot(history.history['accuracy'], label='Train Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.title('Model Accuracy')
plt.legend()
plt.show()
predictions = model.predict(x_test[:5])
for i in range(5):
    plt.imshow(x_test[i], cmap='gray')
    plt.title(f"Predicted: {predictions[i].argmax()}")
    plt.axis('off')
    plt.show()
```

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz>

11490434/11490434 — 1s 0us/step

/usr/local/lib/python3.11/dist-packages/keras/src/layers/reshaping/Flatten.py:37: UserWarning: Do not pass an "input_shape"/"input_dim" argument to a layer. When using Sequential models, prefer using an "Input(shape)" object as the super().__init__(**kwargs)

Epoch 1/5

1875/1875 — 11s 5ms/step - accuracy: 0.8780 - loss: 0.4321 - val_accuracy: 0.9606 - val_loss: 0.1322

Epoch 2/5

1875/1875 — 9s 5ms/step - accuracy: 0.9647 - loss: 0.1210 - val_accuracy: 0.9675 - val_loss: 0.1034

Epoch 3/5

1875/1875 — 10s 5ms/step - accuracy: 0.9753 - loss: 0.0772 - val_accuracy: 0.9724 - val_loss: 0.0881

Epoch 4/5


1875/1875 — 11s 5ms/step - accuracy: 0.9833 - loss: 0.0553 - val_accuracy: 0.9754 - val_loss: 0.0769

Epoch 5/5

1875/1875 — 9s 5ms/step - accuracy: 0.9881 - loss: 0.0419 - val_accuracy: 0.9732 - val_loss: 0.0675

313/313 — 1s 2ms/step - accuracy: 0.9604 - loss: 0.1032

Test Accuracy: 0.9732



Commands + Code + Text ▶ Run all

Connect

plt.show()

↑ ↓ ↻ 🔍 🗨 ⚙ 📄 🗑 ⋮

Predicted: 7



Predicted: 2

