

9/10/25

Week-8

Aim: To implement an LSTM model for predicting sequential data and analyze its performance on a time-series dataset

objective:-

- 1.) understand the working of LSTM Network.
- 2.) Train an LSTM model on sequential time series data
- 3.) Evaluate prediction accuracy and observe the learning behaviour over epochs.

Algorithm:-

→ Data Preprocessing:-

Normalize data and split into train / test sets.

→ Model Design:-

Create LSTM layers with Input, hidden and output layers.

→ Training:-

Feed sequences into LSTM, compute loss and optimize weights using back propagation through time.

→ Prediction:-

Use the trained model to predict the future values.

→ Evaluation:-

Compare predictions with actual values using metrics like MSE (or) RMSE.

Pseudo Code:-

Load dataset

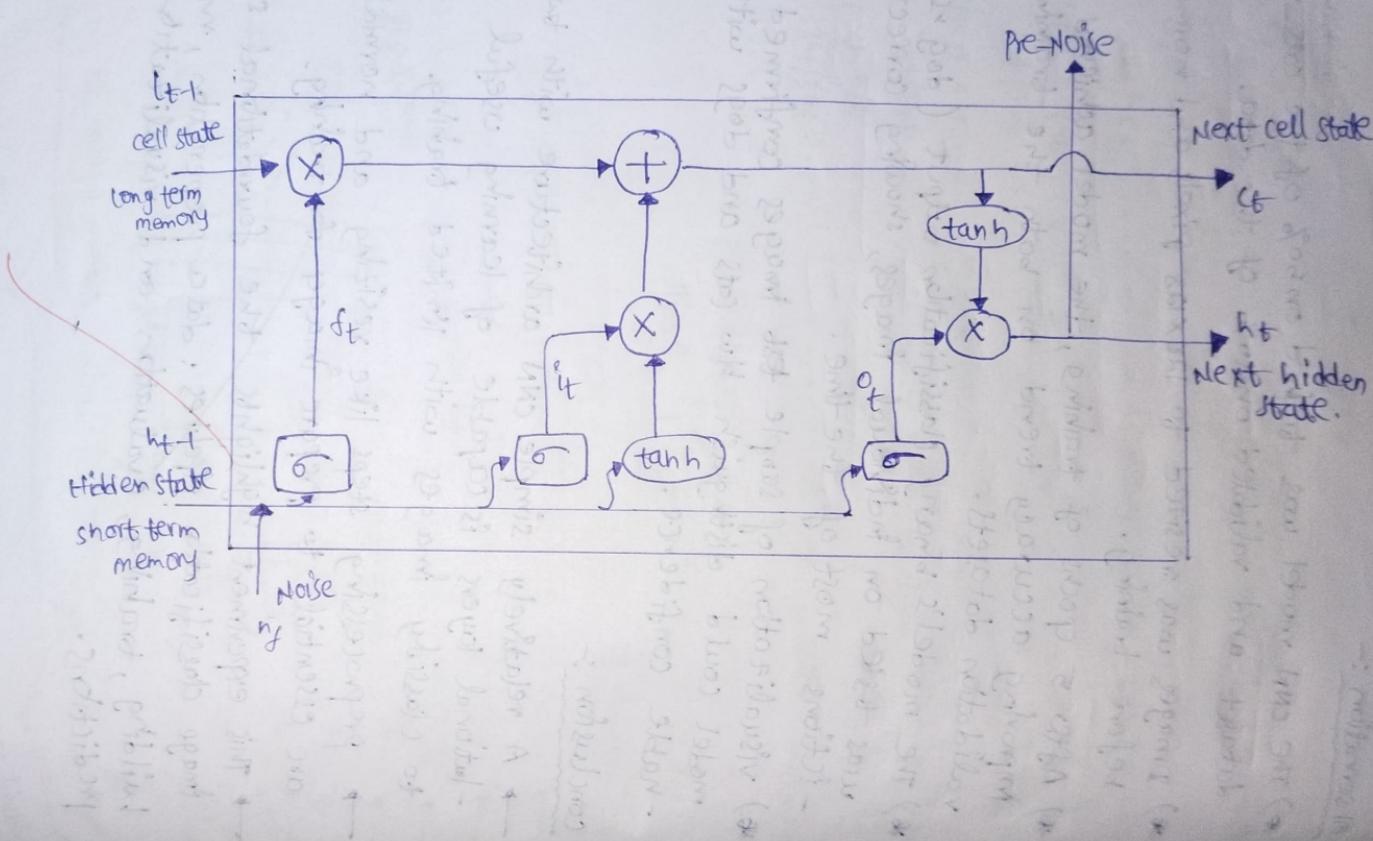
Normalize data

split dataset into train and test

Initialize LSTM model

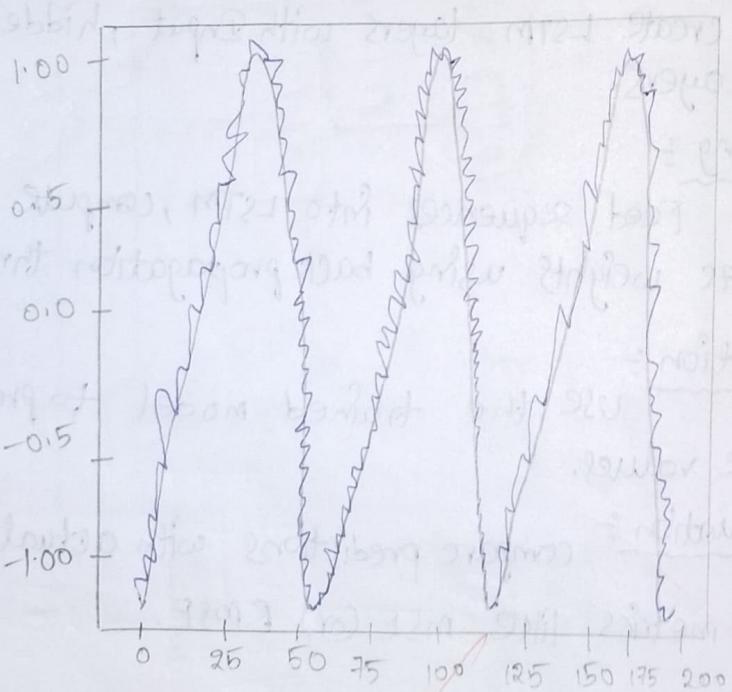
for each epoch

The Structure of LSTM model



Output :

<u>Epoch</u>	<u>Step-loss</u>	<u>val-loss</u>
1	0.2998	0.1087
2	0.0705	0.0050
3	0.0044	0.0011
4	7.9828×10^{-4}	3.5373×10^{-4}
5	3.9283×10^{-4}	3.0203×10^{-4}
6	2.1934×10^{-4}	1.3776×10^{-4}
7	1.6286×10^{-4}	1.0536×10^{-4}



For each batch in training data
predict output
compute loss
back propagate loss
update weights

predict on test data
denormalize predictions
calculate evaluation metrics (MSE / RMSE)

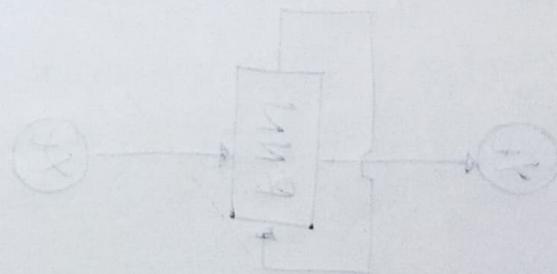
Observation :-

- *) Loss decreases gradually over epochs.
- *) LSTM captures temporal dependencies better than simple RNN's.
- *) Predictions closely follow the trend of actual data but may slightly lag behind sharp changes.

Conclusion :-

- *) LSTM is effective for sequential and time series.
- *) It overcomes the vanishing gradient problems of traditional RNN's.

~~eff. in two~~



```
import tensorflow as tf
from tensorflow.keras.datasets import imdb
from tensorflow.keras.preprocessing.sequence import pad_sequences
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Embedding, LSTM, Dense, Dropout
import matplotlib.pyplot as plt
# Load IMDB dataset
vocab_size = 10000
maxlen = 200
(x_train, y_train), (x_test, y_test) = imdb.load_data(num_words=vocab_size)
x_train = pad_sequences(x_train, maxlen=maxlen)
x_test = pad_sequences(x_test, maxlen=maxlen)
# Build LSTM model
lstm_model = Sequential([
    Embedding(input_dim=vocab_size, output_dim=128, input_length=maxlen),
    LSTM(128, dropout=0.2, recurrent_dropout=0.2),
    Dense(1, activation='sigmoid')
])
lstm_model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
# Train model
lstm_history = lstm_model.fit(
    x_train, y_train,
    epochs=5,
    batch_size=64,
    validation_split=0.2
)
# Evaluate model
lstm_score = lstm_model.evaluate(x_test, y_test)
print(f"LSTM Test Accuracy: {lstm_score[1]*100:.2f}% | Loss: {lstm_score[0]:.4f}")
# Plot Accuracy and Loss
plt.figure(figsize=(12,5))
# Accuracy
plt.subplot(1,2,1)
plt.plot(lstm_history.history['accuracy'], label='Train Accuracy')
plt.plot(lstm_history.history['val_accuracy'], label='Validation Accuracy')
plt.title('LSTM Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()
# Loss
plt.subplot(1,2,2)
plt.plot(lstm_history.history['loss'], label='Train Loss')
plt.plot(lstm_history.history['val_loss'], label='Validation Loss')
plt.title('LSTM Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.show()
```

Downloading data from https://storage.googleapis.com/census-income/tf_keras_dataset/tf_keras_dataset.csv
17464789/17464789 0s 0us/step

Epoch 1/5
/usr/local/lib/python3.12/dist-packages/keras/src/layers/core/embedding.py:97: UserWarning: Argument `input_length` is deprecated. Just remove it.
 warnings.warn(
313/313 193s 603ms/step - accuracy: 0.6676 - loss: 0.5828 - val_accuracy: 0.8122 - val_loss: 0.4151

Epoch 2/5
313/313 198s 593ms/step - accuracy: 0.8517 - loss: 0.3529 - val_accuracy: 0.8358 - val_loss: 0.3951

Epoch 3/5
313/313 206s 606ms/step - accuracy: 0.8840 - loss: 0.2872 - val_accuracy: 0.8328 - val_loss: 0.3862

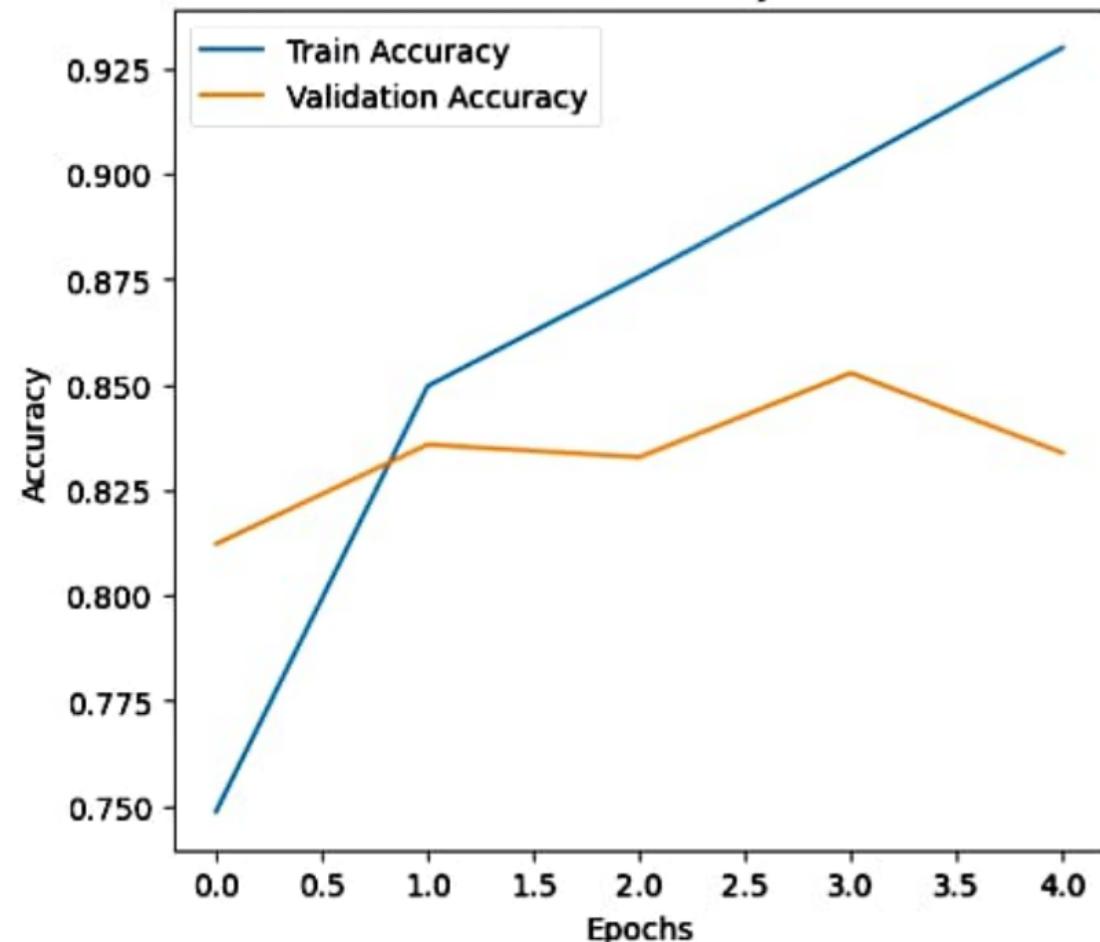
Epoch 4/5
313/313 204s 611ms/step - accuracy: 0.9019 - loss: 0.2525 - val_accuracy: 0.8528 - val_loss: 0.3639

Epoch 5/5
313/313 198s 598ms/step - accuracy: 0.9328 - loss: 0.1832 - val_accuracy: 0.8338 - val_loss: 0.3972

782/782 62s 79ms/step - accuracy: 0.8309 - loss: 0.4013

LSTM Test Accuracy: 83.48% | Loss: 0.3946

LSTM Accuracy



LSTM Loss

